


Convolution Blur Filter in a COMPUTE Shader

Filter

Input: A 2D image, and an $n \times n$ kernel of weights (which sum to one)

Output: A 2D image, where each output pixel is the weighted average of a corresponding $n \times n$ square of input pixels times the respective kernel weights.

What we want: (This is $O(n^2)$)




What we implement: (This is equivalent but only $O(2n)$.)

Write kernel as product of row and column weights

Perform filter in two steps, one horizontal, one vertical

$$\begin{array}{|c|} \hline \text{---} \\ \hline \end{array} = \begin{array}{|c|c|} \hline & \\ \hline \end{array}$$

$N \times 1 * 1 \times N = N \times N$



Building the kernel weights

Build the weight array from a Gaussian bell curve. The true bell curve has infinite tails approaching zero asymptotically. We clamp the range of values to avoid ineffective near-zero weights.

- The width of the kernel is $2w+1$ where the half-width is w .
- The values of the weights are calculated with

$$e^{-\frac{1}{2}\left(\frac{i}{s}\right)^2} \text{ for } i \text{ in range } -w \cdots w$$

where $s=w/2$ controls the width of the bell curve to match the number of desired weights.

- Beware: The range of weights is best considered symmetrically as $-w \cdots w$ whereas the storage in an array must index $0 \cdots 2w$.
- Normalize the array of weights to sum to one.

Efficient filtering in a compute shader

See page 16 of presentation:

<http://amd-dev.wpengine.netdna-cdn.com/wordpress/media/2012/10/Efficient%20Compute%20Shader%20Programming.pps>

Idea

128 threads: (arranged conceptually into a 128×1 row)
will read $128 + 2w$ pixels into shared memory
one pixel each thread, plus one extra for first $2w$ threads
will compute (and write) 128 output pixels
each as a sum of $2w + 1$ weights times $2w + 1$ pixels

Application will create, use, and dispatch (i.e., run) the shader in thread-groups (tiles) which cover the full image:

Tile thread groups (tiles) will be blocks of size 128×1

The dispatch will issue $width/128 \times height$ thread groups.

Shader steps:

Inputs (uniform variables):

src, dst images

w: half-size of kernel

weights: array of $2w + 1$ floats

Declare thread group to be 128×1

Declare thread-group-shared-memory **v[128+2*w+1]** floats

actually must be constant size: **v[128+<largest filter size>]**

Compute **gpos** as the position of the output pixel,

and the center of the $2w + 1$ input pixels

Compute local-index **i** within the thread group

Every thread reads and stores one pixel from src image into shared array

v[i]=imageLoad(src, gpos+ivec2(-w,0))

Some threads (say the first $2w$ of them) load an extra pixel out beyond 128

v[i+128]=imageLoad(src, gpos+ivec2(128-w,0))

Force synchronization between the filling of shared memory (above) and its use (below).

Compute sum of **weights[0 ... 2w]** times corresponding pixels **v[i ... i+2w]**

Store sum at **gpos** in **dst** image

imageStore(dst, gpos, sum)

Compute Shaders

Application code:

Create compute shader

Same as other shaders,
but use **GL_COMPUTE_SHADER** in **glCreateShader** call
Cannot coexist with other shaders in a shader program

CPU invokes computer shader enough times to tile an image:

```
glUseProgram(programID)
// Set all uniform and image variables
glDispatchCompute(W/128, H, 1) // Tiles WxH image with groups sized 128x1
glUseProgram(0)
```

Send block of weights to shader (as a uniform block)

```
glGenBuffers(1, &blockID) // Generates block
bindpoint = ?; // Start at zero, increment for other blocks

loc = glGetUniformLocation(programID, "blurKernel")
glUniformBlockBinding(programID, loc, bindpoint)

glBindBuffer(GL_UNIFORM_BUFFER, blockID)
glBindBufferBase(GL_UNIFORM_BUFFER, bindpoint, blockID)
glBufferData(GL_UNIFORM_BUFFER, #bytes, data, GL_STATIC_DRAW)
```

Send two textures (input and output) to the shader as an image2Ds

```
imageUnit = ?; // Perhaps 0 for input image and 1 for output image

loc = glGetUniformLocation(programID, "...name...") // Perhaps "src" and "dst".
glBindImageTexture(imageUnit, textureID, 0, GL_FALSE, 0, GL_READ_ONLY, GL_RGBA32F)
glUniform1i(loc, imageunit)
// Change GL_READ_ONLY to GL_WRITE_ONLY for output image
// Note: GL_RGBA32F means 4 channels (RGBA) of 32 bit floats.
```

Shader code

```
#version 430 // Version of OpenGL with COMPUTE shader support
layout (local_size_x = 128, local_size_y = 1, local_size_z = 1) in; // Declares thread group size

uniform blurKernel {float weights[101]; }; // Declares a uniform block

layout (rgba32f) uniform readonly image2D src; // src image as 4 channel 32bit float readonly
layout (rgba32f) uniform writeonly image2D dst; // dst image as 4 channel 32bit float writeonly

shared float v[128+101]; // Variable shared with other threads in the 128x1 thread group

void main() {
    ...
    ivec2 gpos = ivec2(gl_GlobalInvocationID.xy); // Combo of groupID, groupSize and localID

    uint i = gl_LocalInvocationID.x; // Local thread id in the 128x1 thread groups128x1

    v[i] = imageLoad(src, gpos+...); // read an image pixel at an ivec2(..) position
    if (i<2*w) v[i+128] = imageLoad(src, gpos+...); // read extra 2*w pixels

    barrier(); // Wait for all threads to catch up before reading v[]
    ...

    imageStore(dst, gpos, ...); // Write to destination image
```